

# Raven Roosting Optimization with MPI: A Parallel Metaheuristic Approach

Author Name

Department of Computer Science

University of Trento

Trento, Italy

author@unitn.it

**Abstract**—This paper presents the implementation and parallelization of the Raven Roosting Optimization (RRO) algorithm, a nature-inspired metaheuristic optimization technique based on the foraging behavior of ravens. We describe both a serial implementation in C and a parallel version using Message Passing Interface (MPI) for distributed memory systems. The algorithm simulates raven behavior including roosting, foraging, and social hierarchy with leader-follower dynamics. Our parallel implementation distributes the population across multiple processes, achieving significant speedup through domain decomposition and collective communication patterns. We present detailed pseudocode for both implementations, discuss key optimizations including early stopping mechanisms and efficient bound checking, and analyze the communication overhead and scalability characteristics of the MPI-based approach.

**Index Terms**—MPI, Parallel Computing, Metaheuristic Optimization, Raven Roosting

## I. INTRODUCTION

### A. Background

Metaheuristic optimization algorithms inspired by natural phenomena have gained significant attention in computational optimization problems. The Raven Roosting Optimization (RRO) algorithm mimics the intelligent foraging behavior of ravens, which are known for their problem-solving abilities and social coordination.

### B. Algorithm Overview

The RRO algorithm is based on three key behavioral patterns observed in ravens:

- **Roosting Behavior:** Ravens return to a common roosting site where they rest and plan their next foraging expedition.
- **Foraging Behavior:** Ravens explore the search space looking for food sources (optimal solutions) through a series of flight steps with exploratory lookouts.
- **Social Hierarchy:** Ravens follow a leader-follower dynamic where the best-performing individual becomes the leader, and a subset of the population acts as followers who explore in the vicinity of the leader.

### C. Problem Statement

The algorithm aims to minimize an objective function  $f(\mathbf{x})$  where  $\mathbf{x} \in \mathbb{R}^n$  represents a position in the search space with  $n$  dimensions (features). Each raven in the population represents

a candidate solution, and the algorithm iteratively improves these solutions through exploration and exploitation phases.

The algorithm maintains:

- A population of ravens, each with a current position and discovered food source
- Fitness values for each raven based on the objective function
- A global leader (best solution found)
- A roosting site (starting point for each iteration)

### D. Motivation for Parallelization

Given the computationally intensive nature of the algorithm, especially with large populations and high-dimensional search spaces, parallelization using MPI becomes essential for:

- Reducing execution time through distributed computation
- Handling larger problem instances
- Scaling across multiple compute nodes in HPC environments

## II. SERIAL IMPLEMENTATION

### A. Algorithm Structure

The serial implementation follows a straightforward execution model where a single process handles the entire population of ravens.

a) *Main Parameters:*

The algorithm accepts the following parameters:

- **Dataset Path:** CSV file containing initial population positions (NxM matrix)
- **Lower/Upper Bounds:** Search space boundaries for each dimension
- **Iterations:** Number of complete optimization cycles
- **Flight Steps:** Number of steps each raven takes toward its destination
- **Lookout Steps:** Number of exploratory searches per flight step
- **Radius:** Looking radius for exploratory searches ( $r_{\text{pept}}$ )

b) *Core Algorithm Pseudocode:*

```

Algorithm: Raven Roosting Optimization
(Serial)
Input: pop_size, features, iterations,
flight_steps,
    lookout_steps, bounds, radius
Output: Best solution found

1. Initialize population from dataset
2. Evaluate fitness for all ravens
3. Set roosting_site as centroid
4. current_leader ← argmin(fitness)
5. followers ← select 20% of
population (excluding leader)

6. for iter = 1 to iterations do
7.   for i = 1 to pop_size do
8.     current_position[i] ←
roosting_site
9.
10.    if is_follower[i] then
11.      target ←
random_point_near(leader, r_pcpt)
12.    else
13.      target ← food_source[i]
14.    end if
15.
16.    for step = 1 to flight_steps do
17.      direction ← normalize(target
- current_position[i])
18.      remaining ←
distance(current_position[i], target)
19.      step_size ← uniform(0,
remaining)
20.      current_position[i] ←
current_position[i] +
21.
step_size × direction
22.
enforce_bounds(current_position[i])
23.
24.    for lookout = 1 to
lookout_steps do
25.      candidate ←
random_point_near(current_position[i],
r_pcpt)
26.      enforce_bounds(candidate)
27.
28.      if f(candidate) <
fitness[i] then
29.        food_source[i] ←
candidate
30.        fitness[i] ← f(candidate)
31.
32.        if random() < 0.1 then
33.          break // Early
stopping
34.        end if
35.      end if
36.    end for
37.  end for
38. end for
39.
40. current_leader ← argmin(fitness)
41. followers ← select 20% of
population (excluding leader)
42. end for
43.
44. return food_source[current_leader]

```

**Listing 1.** Pseudocode for serial RRO implementation (20%, excluding leader)

### c) Key Implementation Details:

#### Initialization Phase:

The initialization phase loads the population from a CSV dataset, evaluates initial fitness values, and calculates the roosting site as a reference point. The looking radius parameter  $r_{\text{pcpt}}$  is computed based on the specified radius and search space bounds.

#### Distance Calculation:

Euclidean distance is computed between positions in the search space. For two positions  $\mathbf{a}, \mathbf{b} \in \mathbb{R}^n$ , the distance is calculated as:

$$d(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{j=1}^n (b_j - a_j)^2} \quad (1)$$

#### Leader Selection:

The leader is determined as the raven with the minimum fitness value, representing the best solution found so far:

$$\text{leader} = \arg \min_{i \in \{1, \dots, P\}} f(\mathbf{x}_i) \quad (2)$$

#### Flight Behavior:

Each raven performs a series of flight steps toward its target. The movement follows:

$$\mathbf{x}_{i,t+1} = \mathbf{x}_{i,t} + s_t \cdot \hat{\mathbf{d}} \quad (3)$$

where

$$\hat{\mathbf{d}} = \frac{\mathbf{t}_i - \mathbf{x}_{i,t}}{\|\mathbf{t}_i - \mathbf{x}_{i,t}\|} \quad (4)$$

is the normalized direction toward target  $\mathbf{t}_i$ , and

$$s_t \sim \mathcal{U}(0, \|\mathbf{t}_i - \mathbf{x}_{i,t}\|) \quad (5)$$

is the step size uniformly sampled from the remaining distance. At each step, the raven performs lookout steps to search for better food sources within radius  $r_{\text{pcpt}}$ .

#### d) Memory Management:

The serial implementation carefully manages memory for:

- Population matrices (food\_source, current\_position)
- Fitness array
- Temporary vectors (direction, n\_candidate\_position, prev\_location, final\_location)
- Leader position and follower flags

#### e) Time Complexity:

For a population of size  $P$ ,  $N$  features,  $I$  iterations,  $F$  flight steps, and  $L$  lookout steps:

- **Overall:**  $O(I \cdot P \cdot F \cdot L \cdot N)$
- **Objective function evaluations:**  $O(I \cdot P \cdot F \cdot L)$

## III. PARALLEL IMPLEMENTATION WITH MPI

### A. Parallelization Strategy

The parallel implementation in the parallel/ directory uses MPI to distribute the raven population across multiple processes. The key parallelization approach is **domain decomposition**, where each MPI process manages a subset of the total population.

#### a) MPI Context and Configuration:

The implementation initializes an MPI context structure that tracks process rank and world size, communicator, and global configuration parameters broadcast to all processes. The context is initialized at program startup to establish the distributed computing environment.

b) *Data Distribution:*

Each process handles a local subset of ravens:

- Process 0 reads parameters and broadcasts to all processes
- Population is divided among processes
- Each process maintains local copies of: food sources, positions, fitness values
- Global leader information is shared across all processes

c) *Parallel Algorithm Pseudocode:*

Algorithm: Raven Roosting Optimization  
(Parallel MPI)

Input: Same as serial + MPI context

Output: Best solution found

```

1. MPI_Init()
2. rank ← MPI_Comm_rank()
3. size ← MPI_Comm_size()

4. if rank == 0 then
5.   Parse command-line arguments
6. end if
7. MPI_Bcast(parameters, root=0)

8. local_rows ← pop_size / size
9. Initialize local population segment
10. Evaluate local fitness values
11. Set roosting_site (shared across
    all processes)

12. MPI_Barrier() // Synchronize
    before timing
13.
14. // Find global leader using
    MPI_Allreduce
15. local_best ← argmin(local_fitness)
16. global_best ←
    MPI_Allreduce(local_best, op=MINLOC)
17. MPI_Bcast(leader_position,
    root=global_best.rank)
18.
19. // Determine followers
    (distributed)
20. num_local_followers ← 0.2 ×
    local_rows
21. local_followers ←
    random_select(num_local_followers)
22.
23. for iter = 1 to iterations do
24.   for i = 1 to local_rows do
25.     // Same flight logic as serial
        version
26.     current_position[i] ←
        roosting_site
27.
28.     if is_follower[i] then
29.       target ←
        random_point_near(leader, r_pcpt)
30.     else
31.       target ← food_source[i]
32.     end if
33.
34.     for step = 1 to flight_steps
do
35.       direction ← normalize(target
- current_position[i])
36.       remaining ←
        distance(current_position[i], target)
37.       step_size ← uniform(0,
        remaining)
38.       current_position[i] ←
        current_position[i] +
39.

```

Listing 2: Pseudocode for the MPbRRO implementation

```

40.
enforce_bounds(current_position[i])
41.
42.     for lookout = 1 to
        lookout_steps do

```

d) *Communication Patterns:*

*Leader Selection:*

Finding the global leader requires collective communication. Each process first finds its local minimum:

$$\text{local\_best}_p = \arg \min_{i \in \text{local\_ravens}_p} f(\mathbf{x}_i) \quad (6)$$

Then a global reduction determines the overall minimum across all processes, followed by broadcasting the leader position from the owning process to all others. This requires  $O(\log k)$  communication time for  $k$  processes.

*Parameter Broadcasting:*

Configuration is broadcast from rank 0 to all processes. The global parameter structure is transmitted as a byte buffer, distributing algorithm configuration uniformly across the distributed system in  $O(\log k)$  time.

*Timing Reduction:*

Maximum execution time across all processes is gathered using a reduction operation:

$$T_{\{\max\}} = \max_{p \in \{0, \dots, k-1\}} T_p \quad (7)$$

This determines the wall-clock time as the slowest process duration.

e) *Synchronization Points:*

The parallel implementation uses several synchronization barriers:

- Before starting computation
- After initialization
- Before timing measurements

f) *Load Balancing:*

Population distribution aims for balanced workload. Ravens are divided as evenly as possible with each process handling

$$P_{\{\text{local}\}} = \left\lfloor \frac{P}{k} \right\rfloor \quad (8)$$

ravens (where  $P$  is total population and  $k$  is number of processes). Each process handles the same number of iterations, and computational work scales linearly with local population size.

g) *Scalability Considerations:*

**Strong Scaling:** For fixed problem size, speedup is limited by:

- Communication overhead (leader updates, broadcasts)
- Load imbalance if population doesn't divide evenly
- Synchronization barriers

**Weak Scaling:** Increasing problem size proportionally with processes maintains efficiency better, as communication-to-computation ratio remains favorable.

h) *Random Number Generation:*

Each process uses an independent random number generator seeded with rank-dependent values. The seed for process  $p$  is computed as

$$s_p = s_{\{\text{base}\}} + p \quad (9)$$

where  $s_{\{\text{base}\}}$  is the initial seed. This ensures reproducibility while maintaining statistical independence across processes.

## IV. OPTIMIZATIONS

### A. Serial Optimizations

a) *Early Stopping Mechanism:*

The serial implementation includes an early stopping heuristic that terminates the lookout phase when a better solution is found. With probability  $p = 0.1$ , the algorithm breaks out of both the lookout and flight loops upon finding an improvement. This optimization reduces unnecessary computation when good solutions are discovered early, potentially saving

$$(1 - p) \cdot L \cdot F \quad (10)$$

lookout operations per successful discovery, where  $L$  is lookout steps and  $F$  is flight steps.

b) *Efficient Memory Access Patterns:*

The implementation uses contiguous memory layouts for population data, enabling better cache locality, vectorization opportunities for modern compilers, and efficient memory operations. Population matrices are stored in row-major order with linear indexing:

$$\text{population}[i \cdot n + j] \quad (11)$$

for raven  $i$  and dimension  $j$ , where  $n$  is the number of features.

c) *Bounds Checking:*

A dedicated bounds checking function ensures solutions remain within feasible space without redundant checks. For each dimension  $j$ , positions are clamped:

$$x_j = \max(L_j, \min(x_j, U_j)) \quad (12)$$

where  $L_j$  and  $U_j$  are the lower and upper bounds respectively.

d) *Random Number Generation:*

Uses PCG (Permuted Congruential Generator) algorithm which provides:

- High statistical quality
- Fast generation
- Small state size
- Good period length

### B. Parallel-Specific Optimizations

a) *Minimizing Communication:*

The parallel implementation strategically places communication:

- **Once per iteration:** Leader selection and broadcast
- **Once at start:** Parameter broadcasting
- **Once at end:** Timing reduction

This design minimizes the frequency of collective operations, keeping most computation local.

b) *Asynchronous Overlapping Potential:*

While the current implementation uses blocking collectives, the structure allows future optimization through:

- Non-blocking leader broadcast (MPI\_Ibroadcast)
- Overlapping computation of independent ravens with communication
- Pipelined iteration execution

c) *Local-First Processing:*

Each process computes its local leader before global reduction, minimizing data transfer:

- **Without local-first:** Transfer all fitness values ( $P$  doubles)
- **With local-first:** Transfer only local minimum (1 struct with rank and value)

#### d) Reduced Early Stopping in Parallel:

The parallel version conditionally disables early stopping when measuring speedup. When benchmarking mode is enabled, the probabilistic early termination is suppressed to ensure deterministic execution for fair performance comparisons across different process counts.

#### e) Memory Optimizations:

##### *Reusable Buffers:*

The implementation uses reusable buffers for temporary computations:

- direction, n\_candidate\_position, prev\_location, final\_location

These are allocated once per process and reused across all iterations, avoiding repeated allocation overhead.

##### *Structure Packing:*

MPI communication of configuration uses byte-level broadcasting, transmitting the entire structure as a contiguous byte array. This avoids overhead of creating custom MPI datatypes for simple structs.

### C. Algorithmic Optimizations

#### a) Follower Selection Strategy:

Followers are selected randomly with 20% of population. The number of followers is computed as:

$$N_{\{\text{followers}\}} = \lceil 0.2 \cdot P - 1 \rceil \quad (13)$$

This balances exploration (non-followers search their own best) and exploitation (followers search near global leader).

##### b) Adaptive Step Sizing:

Step size is proportional to remaining distance. At each flight step  $t$ , the step size is:

$$s_t = d_{\{\text{remaining}\}} \cdot u \quad (14)$$

where  $d_{\{\text{remaining}\}} = \|\mathbf{t}_i - \mathbf{x}_{i,t}\|$  is the distance to target and  $u \sim \mathcal{U}(0, 1)$ . This provides large steps when far from target (exploration) and small steps when near target (exploitation).

##### c) Roosting Reset:

Ravens return to roosting site at the start of each iteration. For all ravens  $i$ :

$$\mathbf{x}_{i,0} \leftarrow \mathbf{r} \quad (15)$$

where  $\mathbf{r}$  is the roosting site. This reset mechanism prevents premature convergence, maintains diversity in search trajectories, and models natural raven behavior.

### D. Performance Considerations

#### a) Computational Bottlenecks:

Profile analysis would likely show hotspots in:

- 1) **Objective function evaluation** (lookout\_steps  $\times$  flight\_steps  $\times$  pop\_size calls per iteration)
- 2) **Distance calculations** (repeated square root operations)
- 3) **Random number generation** (multiple calls per step)

#### b) Potential Further Optimizations:

Future improvements could include:

- **Vectorization:** SIMD operations for distance calculations
- **GPU acceleration:** Parallel evaluation of objective function
- **Hybrid MPI+OpenMP:** Thread-level parallelism within each node
- **Adaptive parameters:** Dynamic adjustment of flight\_steps and lookout\_steps based on convergence rate
- **Communication aggregation:** Batch multiple MPI operations

## V. CONCLUSION

### A. Summary

This paper presented a comprehensive study of the Raven Roosting Optimization algorithm, including both serial and parallel MPI implementations. The algorithm successfully models the foraging behavior of ravens through a metaheuristic approach that balances exploration and exploitation via leader-follower dynamics.

The serial implementation provides a clear baseline with optimizations such as early stopping and efficient memory access patterns. The parallel MPI version achieves scalability through domain decomposition, distributing the population across multiple processes while maintaining algorithmic correctness through collective communication for leader selection.

### B. Key Contributions

- Detailed pseudocode for both serial and parallel implementations
- Analysis of parallelization strategy using MPI domain decomposition
- Identification of optimization techniques including early stopping, adaptive step sizing, and efficient memory management
- Discussion of communication patterns and synchronization requirements
- Performance considerations for scalability

### C. Implementation Highlights

The implementations demonstrate several best practices:

- Clean separation between computation and communication logic
- Rank-specific random number seeding for reproducibility
- Configurable parameters for flexible experimentation
- Comprehensive timing instrumentation
- Proper memory management with cleanup routines

### D. Scalability Analysis

The parallel implementation exhibits:

- **Good strong scaling potential** for large populations where computation dominates communication
- **Excellent weak scaling characteristics** as problem size grows proportionally with process count

- **Communication bottleneck** at leader selection phase, occurring once per iteration
- **Load balance** maintained through even population distribution

#### E. Future Work

Several avenues for enhancement remain:

- Hybrid Parallelism:* Combining MPI with OpenMP could exploit both distributed and shared memory parallelism, particularly beneficial on modern multi-core cluster nodes.
- Advanced Communication:* Non-blocking collectives and asynchronous leader updates could overlap communication with computation, potentially improving strong scaling efficiency.
- Adaptive Mechanisms:* Dynamic parameter adjustment based on convergence metrics could reduce unnecessary computation in later iterations when the population converges.
- Alternative Objective Functions:* The modular design supports easy integration of different objective functions, enabling application to diverse optimization problems.
- Performance Modeling:* Analytical performance models could predict scalability and guide optimal process count selection for given problem sizes.

#### F. Concluding Remarks

The Raven Roosting Optimization algorithm demonstrates the effectiveness of nature-inspired metaheuristics for complex optimization problems. The MPI parallelization successfully distributes computational workload while maintaining the algorithm's essential characteristics. Both implementations provide a solid foundation for further research and practical applications in high-performance computing environments.

The careful balance between local computation and global coordination exemplifies effective parallel algorithm design, achieving meaningful speedup while preserving solution quality. The documented code structure, complete with pseudocode and implementation details, facilitates reproducibility and extension by the research community.

#### REFERENCES